



Process Description and Control

**Week 3
(part one)**

COMP 4735
Rafael Roman Otero

What's an OS?
(small detour)

What's an OS?

“a **program** that controls the execution of application programs, and acts interface between applications and the computer hardware”

--- Stallings

What's an OS?

“a **program** that controls the execution of application programs, and acts interface between applications and the computer hardware”

--- Stallings

“It is hard to pin down what an operating system is other than saying **it is the software that runs in kernel mode**—and even that is not always true.

--- Tanenbaum

What's an OS?

The OS as an extended machine (Tanenbaum)

- The architecture of most computers at the machine-language level is **primitive** and **awkward to program**, especially for I/O.

What's an OS?


The OS as an extended machine (Tanenbaum)

- Say you want to do bare-metal **USB programming**.

Documentation needed:

- USB protocol Book (300 pages) so you can make yourself a semi-expert on USB protocol.
- RC-508 Datasheet (a few hundred pages as well)
- The manual of the chip that is mounted on the RC-508 (the one doing low-level USB stuff)
- A book on PCI protocol (300 pages), or else how will you write the PCI drivers that you need to talk to the device?
- Your computer's chipset manual (100 pages)
- X86 developers manual (2k pages)

(+ 10 months of trying to get the thing to do something)




RC-508

What's an OS?

The OS as an extended machine (Tanenbaum)

- Say you want to do bare-metal **USB programming**.



RC-508

It's ridiculous how complex low-level I/O programming is.

You can't possibly expect a computer to be useful at all when you have to do everything yourself


What's an OS?

The OS as an extended machine (Tanenbaum)

- Say you want to do bare-metal **USB programming**.

Instead an OS provides you and the whole world with:

- a) **PCI Drivers**
- b) **USB Drivers**
- c) Additional **abstractions** like “files”, “drives”, “plug and play” so you can print the complain letter you wrote to the Queen and continue with your life.



RC-508

What's an OS?

The OS as an extended machine (Tanenbaum)

- **In general, CPUs, memories, disks, and I/O devices are very complex and present difficult, awkward, idiosyncratic, and inconsistent interfaces.**

What's an OS?

The OS as an extended machine (Tanenbaum)

- In general, CPUs, memories, disks, and I/O devices are very **complex** and present **difficult, awkward, idiosyncratic, and inconsistent interfaces**.

It's the OS task to hide the hardware and present programs (and their programmers) with **nice, clean, elegant, consistent, abstractions to work with instead**.

What's an OS?

The OS as an extended machine (Tanenbaum)

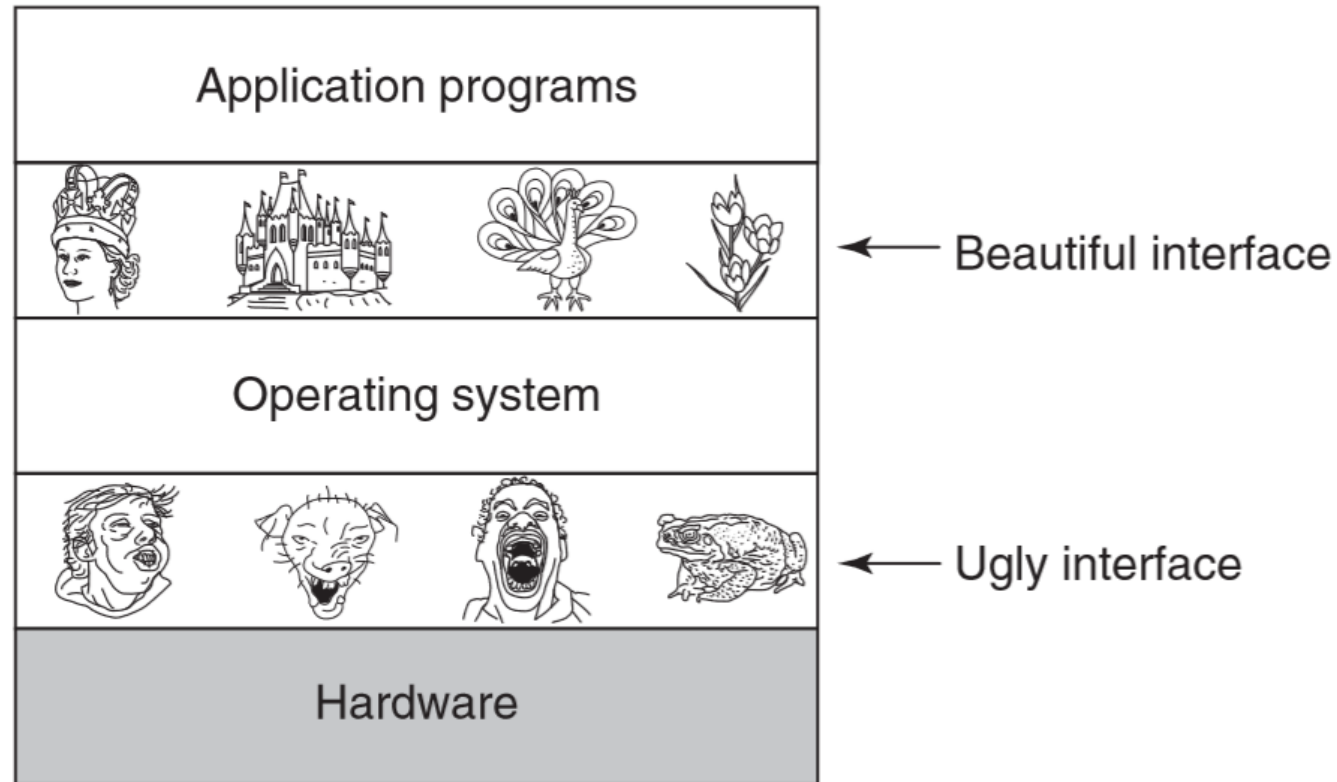


Figure 1-2. Operating systems turn ugly hardware into beautiful abstractions.

What's an OS?

The OS as a resource manager (Tanenbaum)

- An alternative, bottom-up, view holds that the operating system is there to **manage all the pieces of a complex system.**

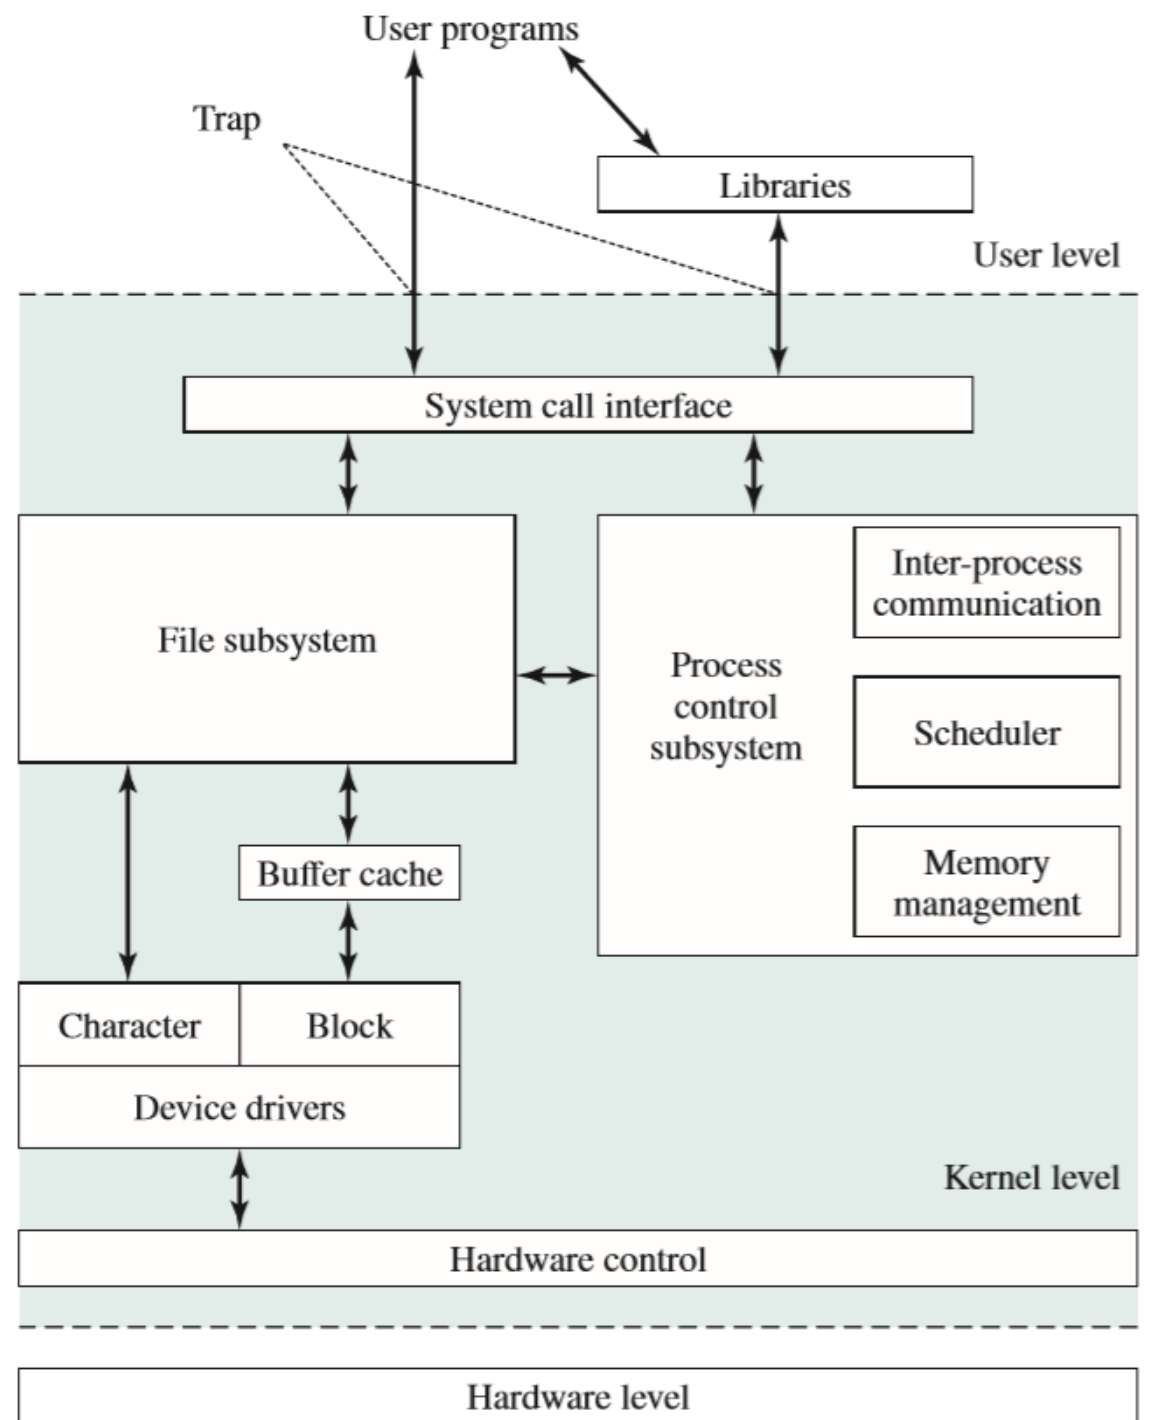
What's an OS?

The OS as a resource manager (Tanenbaum)

- An alternative, bottom-up, view holds that the operating system is there to **manage all the pieces of a complex system**.
- In this view, the OS's job is to provide for an **orderly and controlled allocation** of the processors, memories, and I/O devices among processes.

What's an OS?

Traditional *nix architecture



Stalling's

What's an OS?

Resources on the Web

- [The mind behind Linux | Linus Torvalds](#)
- [Old School Sean - The MINIX operating system](#)

What's an OS?

Book reference
Section 2.1, 2.8

Process

Process

What is a process?

Process

“A program in execution”

“An instance of a program running on a computer”

--- **Stallings**

“A process is basically a program in execution”

--- **Tanenbaum**

Process

Definition revolves around **the task**,
not the program.



“The agent that carries out the task described in a program is called a process.”

--- **Halдар & Aravind**

“An executable that was brought into ‘existence’ by the OS”

--- **My definition**

Process

Each process is associated with an **address space**.



What's an address space?

Process

Each process is associated with an **address space**. This address space contains:

- The process' _____ ?
- The process' _____ ?
- The process' _____ ?

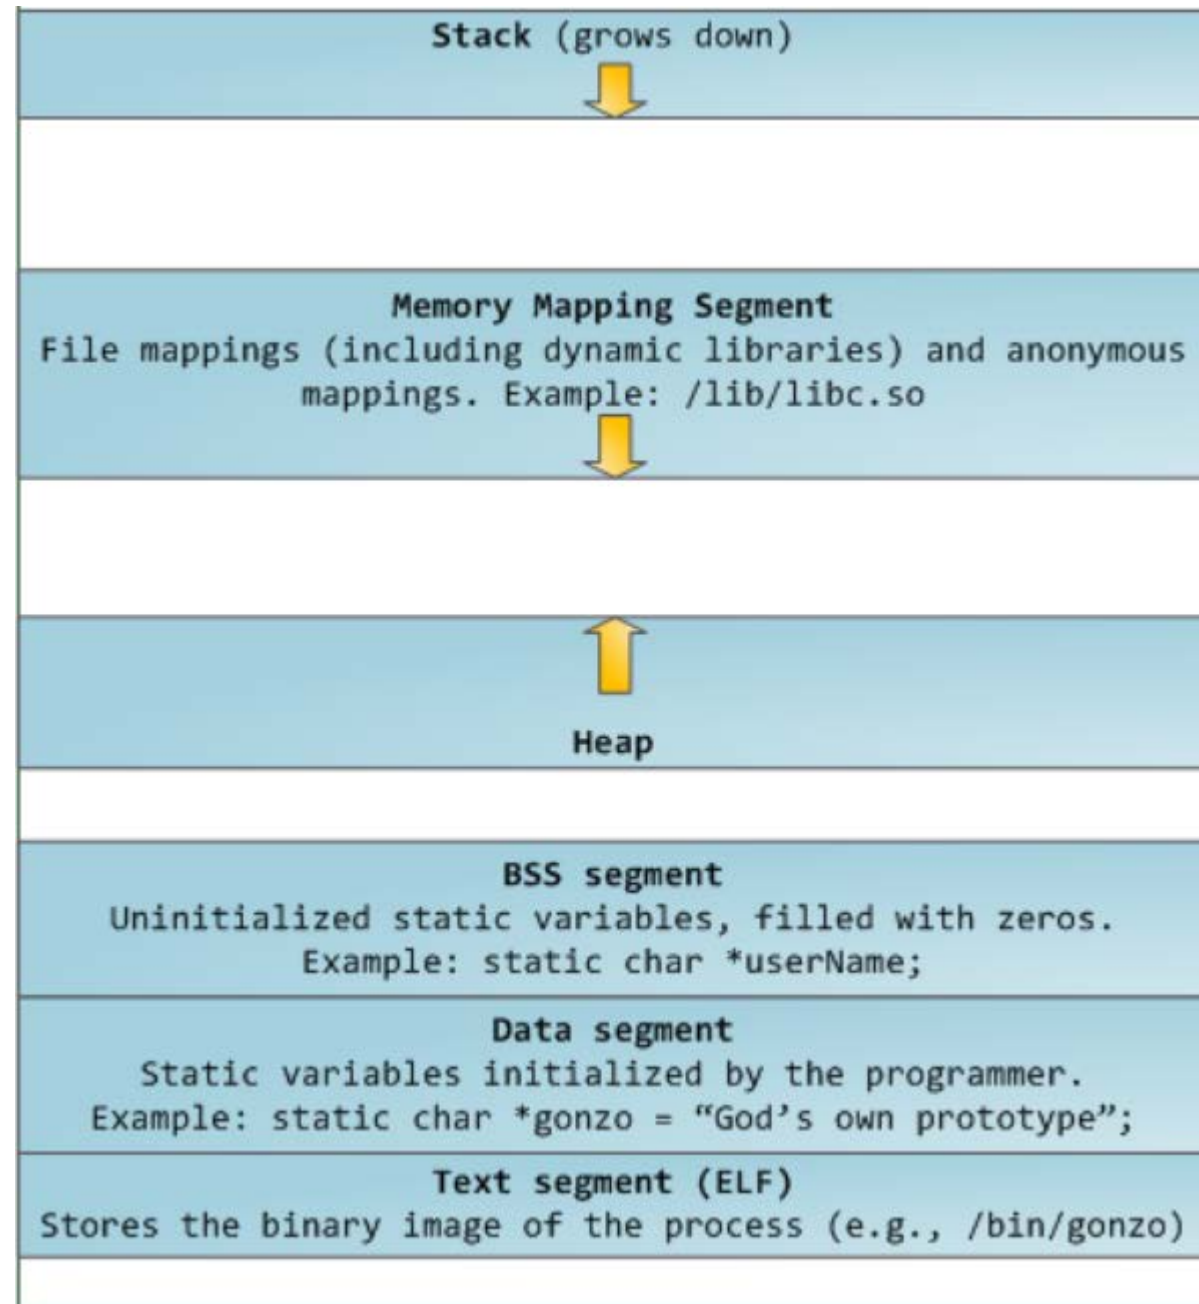
Process

Each process is associated with an **address space**. This address space contains:

- The process' _____ ?
- The process' _____ ?
- The process' _____ ?

Process

- A process address' space (in more detail)



Up to some number
(e.g. 0xFFFFFFFF)

From some number
(e.g. 0x00000000)

ubuntu@ip-172-31-13-223:~\$ htop

htop is a system monitor for Linux

CPU[0.0%] Tasks: 29, 24 thr; 1 running
 Mem[89.9M/984M] Load average: 0.00 0.00 0.00
 Swp[0K/0K] Uptime: 01:03:13

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
1626	ubuntu	20	0	32324	4320	3616	R	0.7	0.4	0:00.02	htop
1	root	20	0	156M	9088	6796	S	0.0	0.9	0:02.33	/sbin/init
401	root	19	-1	94844	12256	11568	S	0.0	1.2	0:00.57	/lib/systemd/systemd-journald
415	root	20	0	97708	1948	1772	S	0.0	0.2	0:00.00	/sbin/lvmtool -f
427	root	20	0	43560	4644	3184	S	0.0	0.5	0:00.26	/lib/systemd/systemd-udev
483	systemd-t	20	0	138M	3236	2720	S	0.0	0.3	0:00.00	/lib/systemd/systemd-timesyncd
454	systemd-t	20	0	138M	3236	2720	S	0.0	0.3	0:00.01	/lib/systemd/systemd-timesyncd
570	systemd-n	20	0	80016	5340	4748	S	0.0	0.5	0:00.01	/lib/systemd/systemd-networkd
588	systemd-r	20	0	70612	5444	4892	S	0.0	0.5	0:00.02	/lib/systemd/systemd-resolved
711	root	20	0	25376	292	60	S	0.0	0.0	0:00.06	/sbin/iscsid
712	root	10	-10	25880	5272	4044	S	0.0	0.5	0:00.00	/sbin/iscsid
727	daemon	20	0	28332	2544	2332	S	0.0	0.3	0:00.00	/usr/sbin/atd -f
735	root	20	0	31748	3232	2944	S	0.0	0.3	0:00.00	/usr/sbin/cron -f

Process

Each process is also associated with a **set of resources**:

- **Registers**
- **File descriptors (used to access files, IO, pipes, sockets)**
- **CPU number, etc**

Process

Processes exists in the context of **multiprogramming** or **multitasking** operating systems.



What's this?

Process

Processes exist in the context of **multiprogramming** or **multitasking** operating systems.

How can several processes be in memory simultaneously?

Task Manager

File Options View

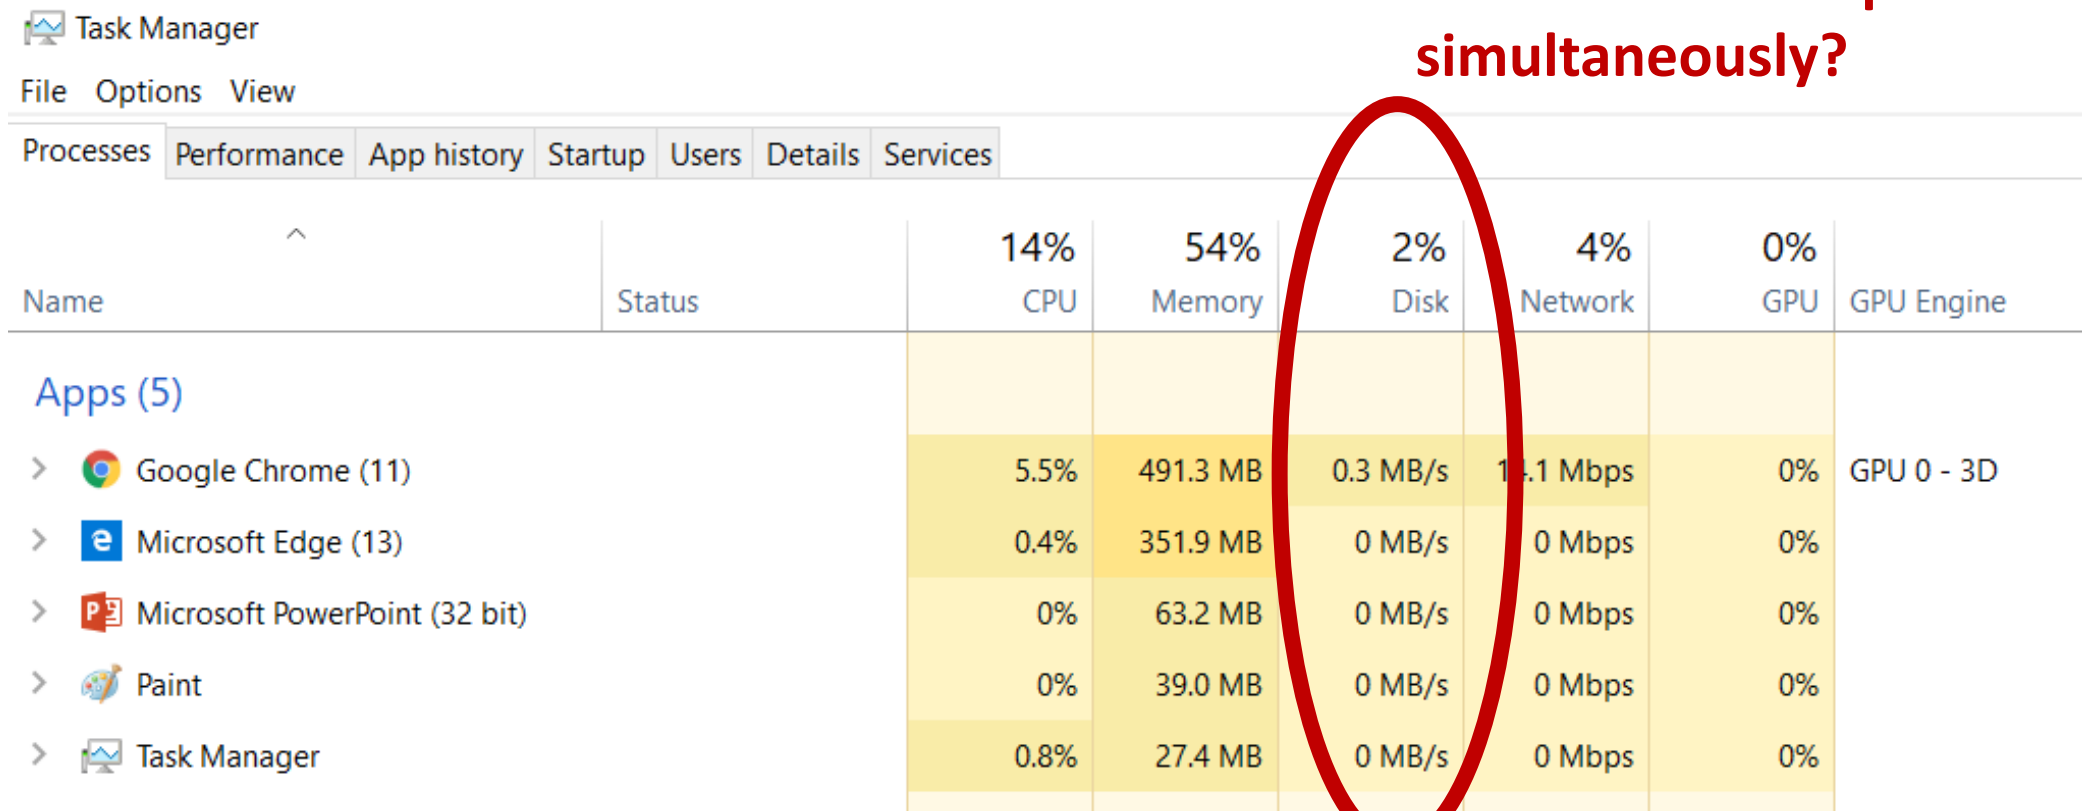
Processes Performance App history Startup Users Details Services

Name	Status	14% CPU	54% Memory	2% Disk	4% Network	0% GPU	GPU Engine
Apps (5)							
> Google Chrome (11)		5.5%	491.3 MB	0.3 MB/s	14.1 Mbps	0%	GPU 0 - 3D
> Microsoft Edge (13)		0.4%	351.9 MB	0 MB/s	0 Mbps	0%	
> Microsoft PowerPoint (32 bit)		0%	63.2 MB	0 MB/s	0 Mbps	0%	
> Paint		0%	39.0 MB	0 MB/s	0 Mbps	0%	
> Task Manager		0.8%	27.4 MB	0 MB/s	0 Mbps	0%	

Process

Processes exist in the context of **multiprogramming** or **multitasking** operating systems.

How can several processes use disk simultaneously?



The screenshot shows the Windows Task Manager Performance tab. The 'Disk' column is highlighted with a red oval, illustrating how multiple processes can use the disk simultaneously. The table below represents the data shown in the screenshot.

Task Manager		Performance					
File Options View		Processes Performance App history Startup Users Details Services					
Name	Status	14% CPU	54% Memory	2% Disk	4% Network	0% GPU	GPU Engine
Apps (5)							
> Google Chrome (11)		5.5%	491.3 MB	0.3 MB/s	14.1 Mbps	0%	GPU 0 - 3D
> Microsoft Edge (13)		0.4%	351.9 MB	0 MB/s	0 Mbps	0%	
> Microsoft PowerPoint (32 bit)		0%	63.2 MB	0 MB/s	0 Mbps	0%	
> Paint		0%	39.0 MB	0 MB/s	0 Mbps	0%	
> Task Manager		0.8%	27.4 MB	0 MB/s	0 Mbps	0%	

Process

Processes exist in the context of **multiprogramming** or **multitasking** operating systems.

Task Manager

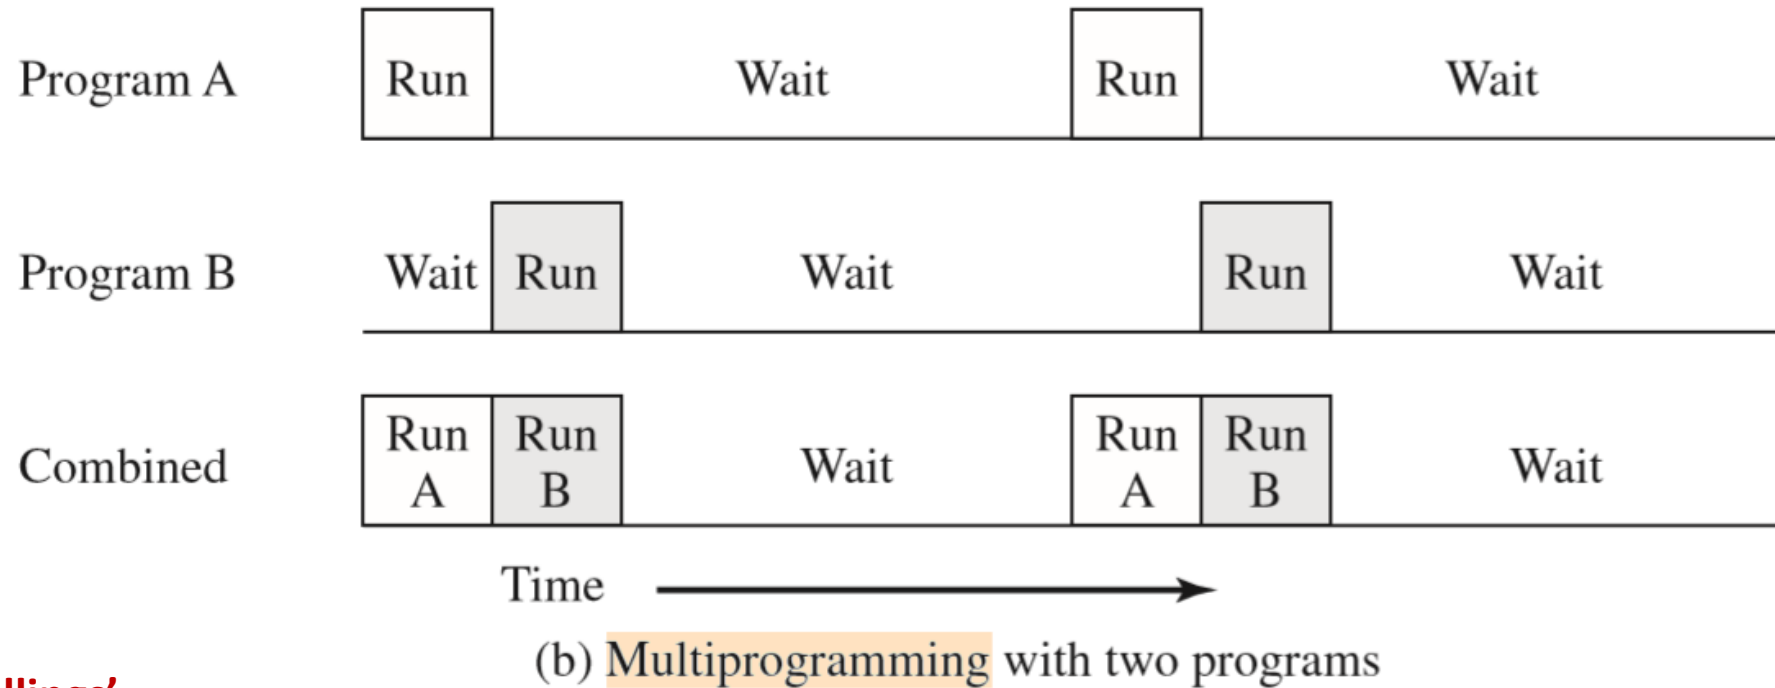

File Options View

Processes Performance App history Startup Users Details Services

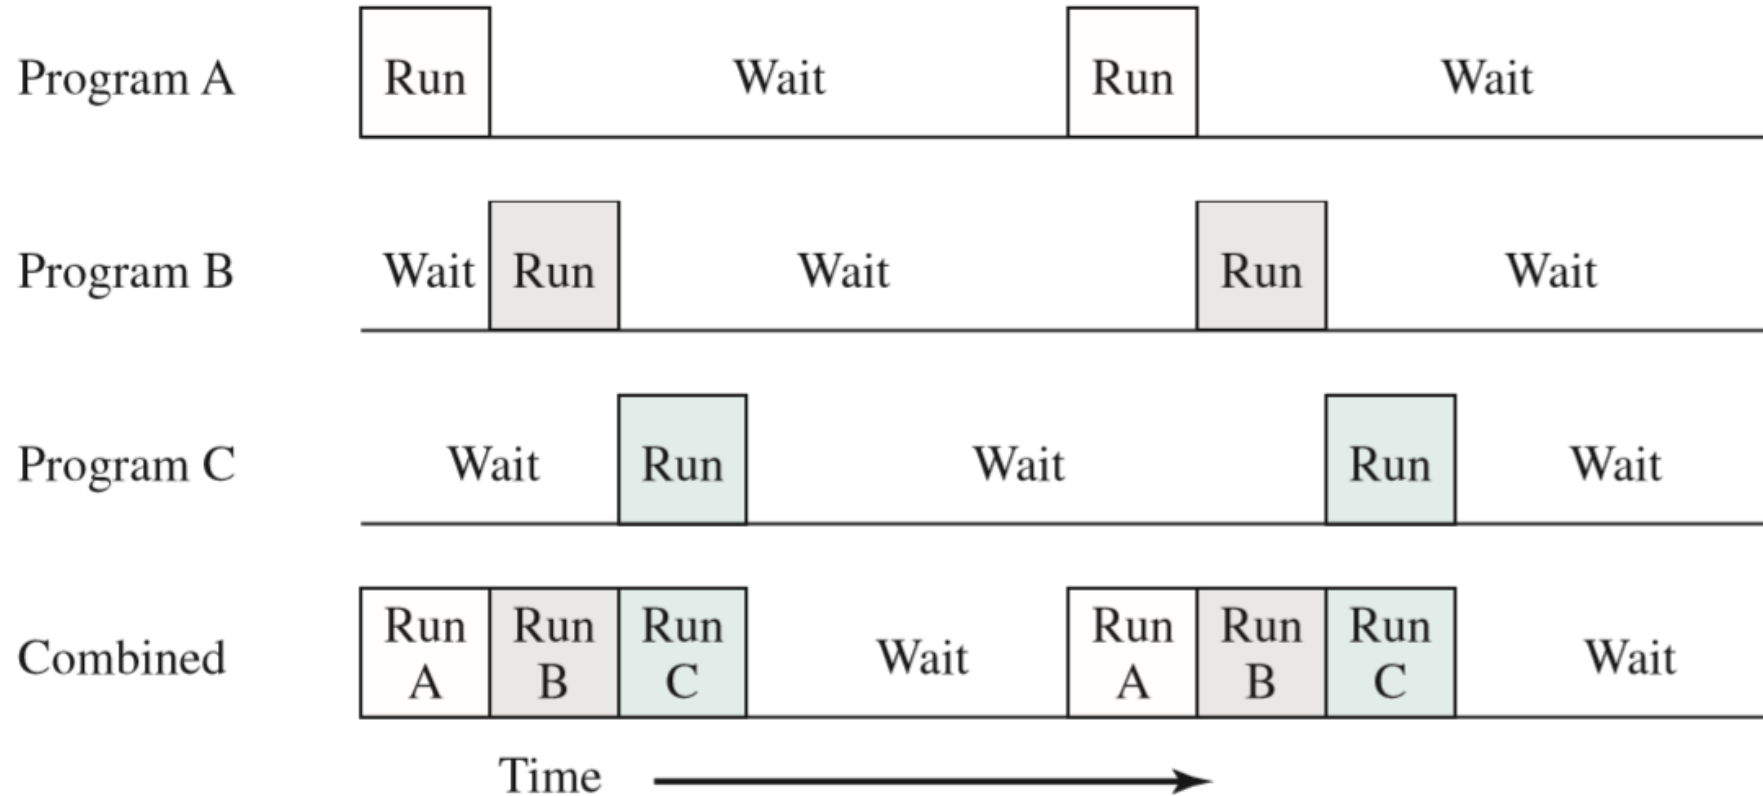
Name	Status	14% CPU	54% Memory	2% Disk	4% Network	0% GPU	GPU Engine
Apps (5)							
> Google Chrome (11)		5.5%	491.3 MB	0.3 MB/s	14.1 Mbps	0%	GPU 0 - 3D
> Microsoft Edge (13)		0.4%	351.9 MB	0 MB/s	0 Mbps	0%	
> Microsoft PowerPoint (32 bit)		0%	63.2 MB	0 MB/s	0 Mbps	0%	
> Paint		0%	39.0 MB	0 MB/s	0 Mbps	0%	
> Task Manager		0.8%	27.4 MB	0 MB/s	0 Mbps	0%	

How can several processes use CPU simultaneously? What does 0% mean?

Process



Process



(c) Multiprogramming with three programs

Process

- In a multiprogramming system, the kernel periodically grants **short periods of CPU access** to processes.
- When the kernel takes the CPU from a process, we say the process is **suspended**.

Process

- In a multiprogramming system, the kernel periodically grants **short periods of CPU access** to processes.
- When the kernel takes the CPU from a process, we say the process is **suspended**.

How long typically?



When does this happen?
(Hint: there's a few occasions)



Process

Once a process is suspended, it needs **certain information** to resume execution. So the Kernel keeps a **process table** (or **process control block**)

Process

Once a process is suspended, it needs **certain information** to resume execution. So the Kernel keeps a **process table** (or **process control block**)

- ID, name
- Process state (suspended, running, waiting for IO, etc)
- Priority
- Stack Pointer
- Program Counter <- - **Maybe in some implementations**
- File descriptors
- Address space
- Current CPU

torvalds / linux

Watch

6,618

Star

68,296

Fork

24,614

<> Code

Pull requests 248

Projects 0

Insights

Branch: master

linux / include / linux / sched.h

Find file

Copy path

torvalds Merge git://git.kernel.org/pub/scm/linux/kernel/git/davem/net

e874644 4 days ago

597 struct task_struct {

593 #ifdef CONFIG_THREAD_INFO_IN_TASK

594 /*

595 * For reasons of header soup (see current_thread_info())

596 * must be the first element of task_struct.

597 */

598 struct thread_info thread_info;

599 #endif

600 /* -1 unrunnable, 0 runnable, >0 stopped: */

601 volatile long state;

636 int on_rq;

637

638 int prio;

639 int static_prio;

640 int normal_prio;

641 unsigned int rt_priority;

642

609 void

610 atomic_t

611 /* Per task flags (PF_*), defined further below: */

612 unsigned int flags;

613 unsigned int ptrace;

614

615 #ifdef CONFIG_SMP

616 struct llist_node wake_entry;

617 int on_cpu;

618

619 #endif

620 #ifdef CONFIG_THREAD_INFO_IN_TASK

621 /* Current CPU: */

622 unsigned int cpu;

623

Process

Max number of threads in Linux?

Process

Max number of threads in Linux?

man7.org > Linux > [man-pages](#)

0X3ff...f = 1,073 Millions

`/proc/sys/kernel/threads-max` (since Linux 2.3.11)

This file specifies the system-wide limit on the number of threads (tasks) that can be created on the system.

Since Linux 4.1, the value that can be written to **threads-max** is bounded. The minimum value that can be written is 20. The maximum value that can be written is given by the constant **FUTEX_TID_MASK** (0x3fffffff). If a value outside of this range is written to **threads-max**, the error **EINVAL** occurs.

The value written is checked against the available RAM pages. If the thread structures would occupy too much (more than 1/8th) of the available RAM pages, **threads-max** is reduced accordingly.

Process

Resources on the Web

- [Linux: Processes](#)

Process

Book reference

Section 3.1 and 2.2 (for multiprogramming definition)

The Life of a Process (process states)

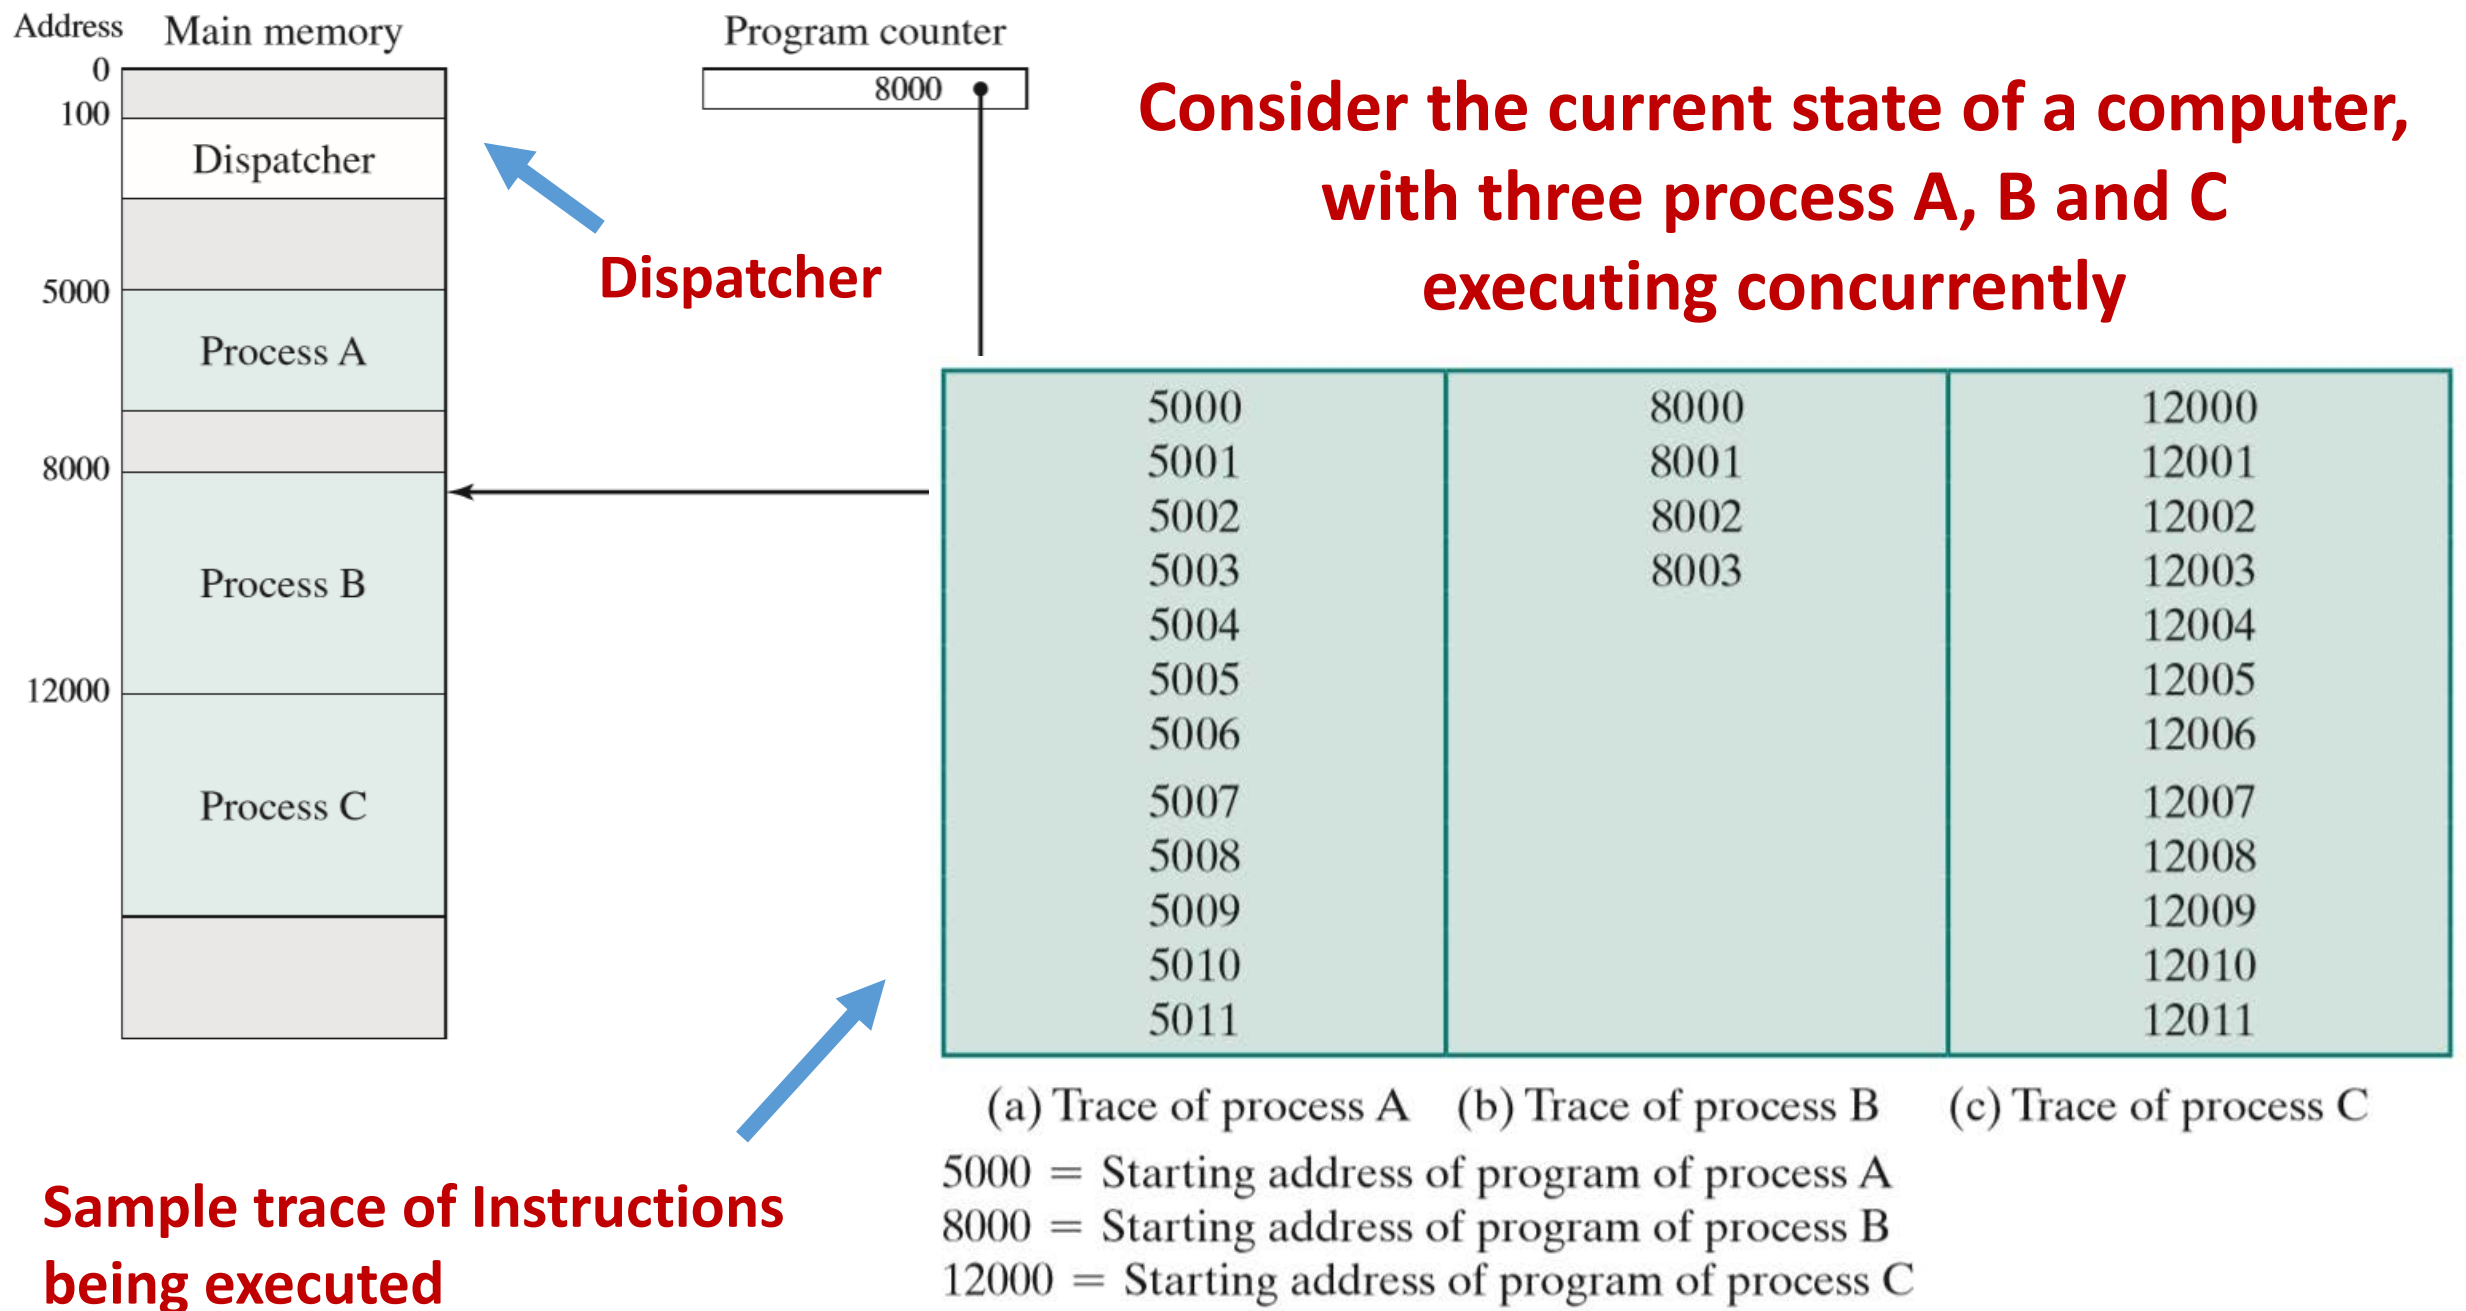


Figure 3.3 Traces of Processes of Figure 3.2

1	5000		27	12004
2	5001		28	12005
3	5002	A		-----Time-out
4	5003		29	100
5	5004		30	101
6	5005		31	102
	-----Time-out		32	103
7	100		33	104
8	101		34	105
9	102	Dispatcher	35	5006
10	103		36	5007
11	104		37	5008
12	105		38	5009
13	8000		39	5010
14	8001	B	40	5011
15	8002			-----Time-out
16	8003		41	100
	-----I/O request		42	101
17	100		43	102
18	101		44	103
19	102	Dispatcher	45	104
20	103		46	105
21	104		47	12006
22	105		48	12007
23	12000		49	12008
24	12001		50	12009
25	12002	C	51	12010
26	12003		52	12011
				-----Time-out

**Sample concurrent execution
for a single core**

Stallings'


The Life of a process

- Newly created processes begin in a **new state**. It's process table was created, but not yet loaded into memory.
(at this point they're **not eligible to get picked by the dispatcher**).



The Life of a process

- Once loaded to memory the move to a ready queue, and their state is now in **ready state**.
(at this point **they are eligible to get picked by the dispatcher**).



The Life of a process


- When their turn comes, they'll be moved to **running state**, and the dispatcher will give them control of the CPU.



The Life of a process


- While in execution, a process terminate execution, thus moving to an **exit state**.

(Processes in exit state have been removed from the ready queue)




The Life of a process

- While in execution, a process could also get placed back to the ready queue and into **ready state**, when a tick occurs.



The Life of a process

- While in execution, a process could also request for IO (say *System.in.readLine()*) and get into a **Blocked state**.
(when blocked, processes wait in the blocked queue)

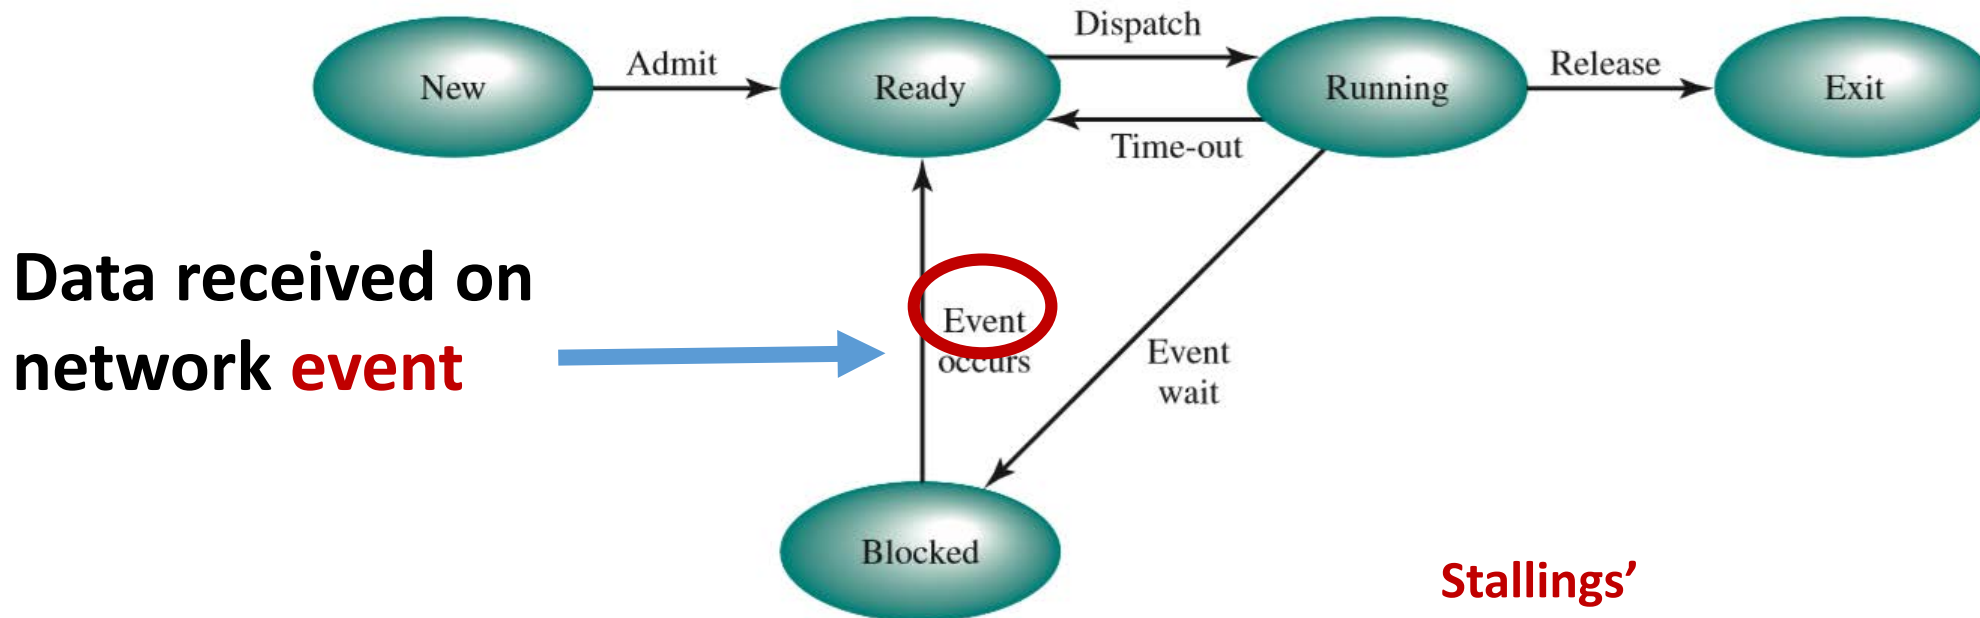


Read data from network
(wait for data received on
network **event**)

Stallings'

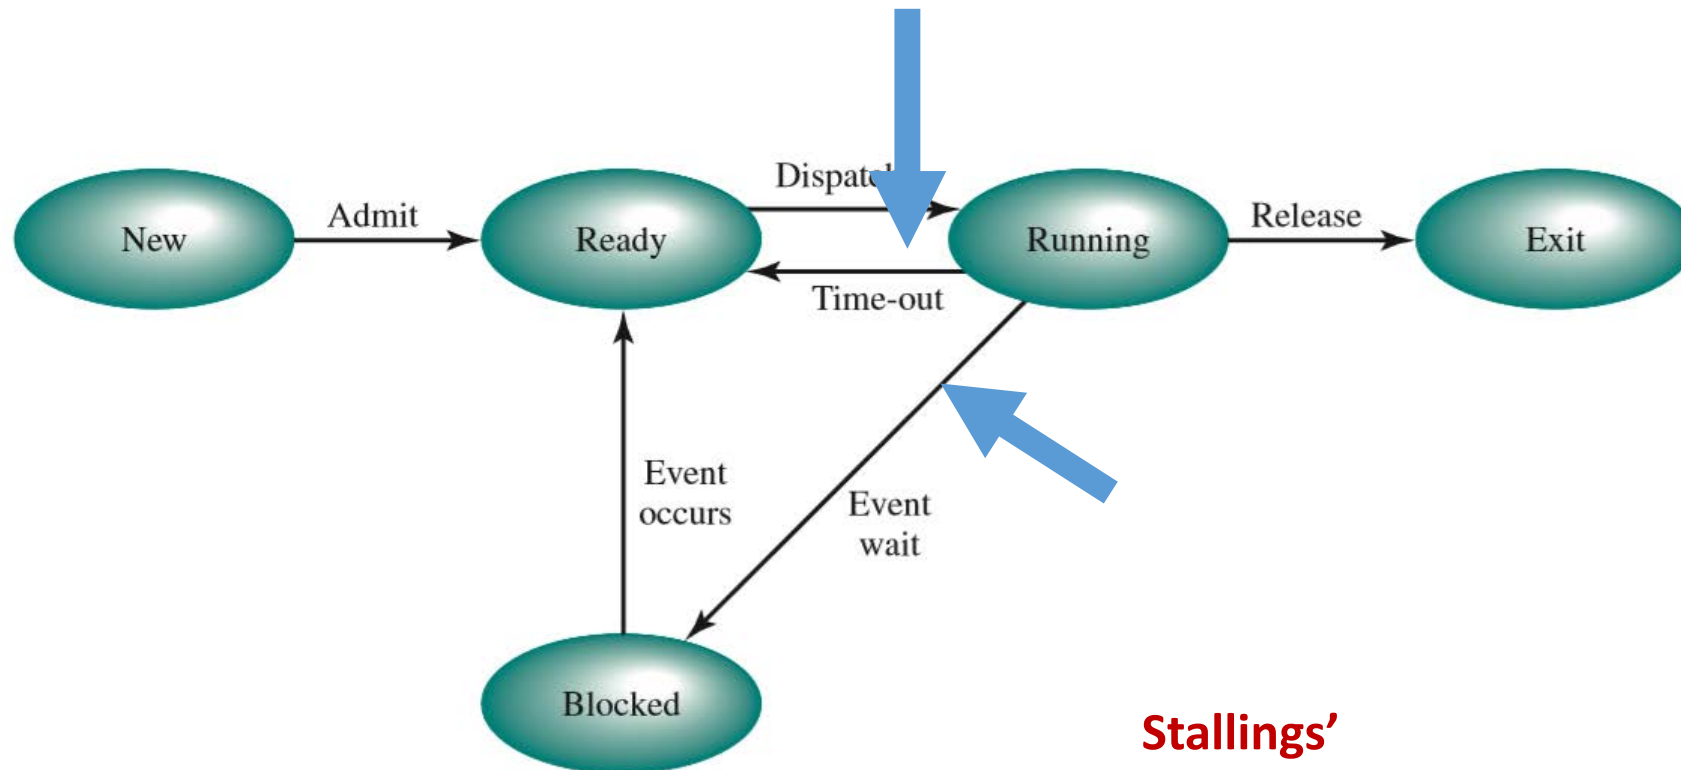
The Life of a process

- A blocked process gets unblocked once the IO it's waiting for arrives, and goes back to **ready state**, and on the ready queue.



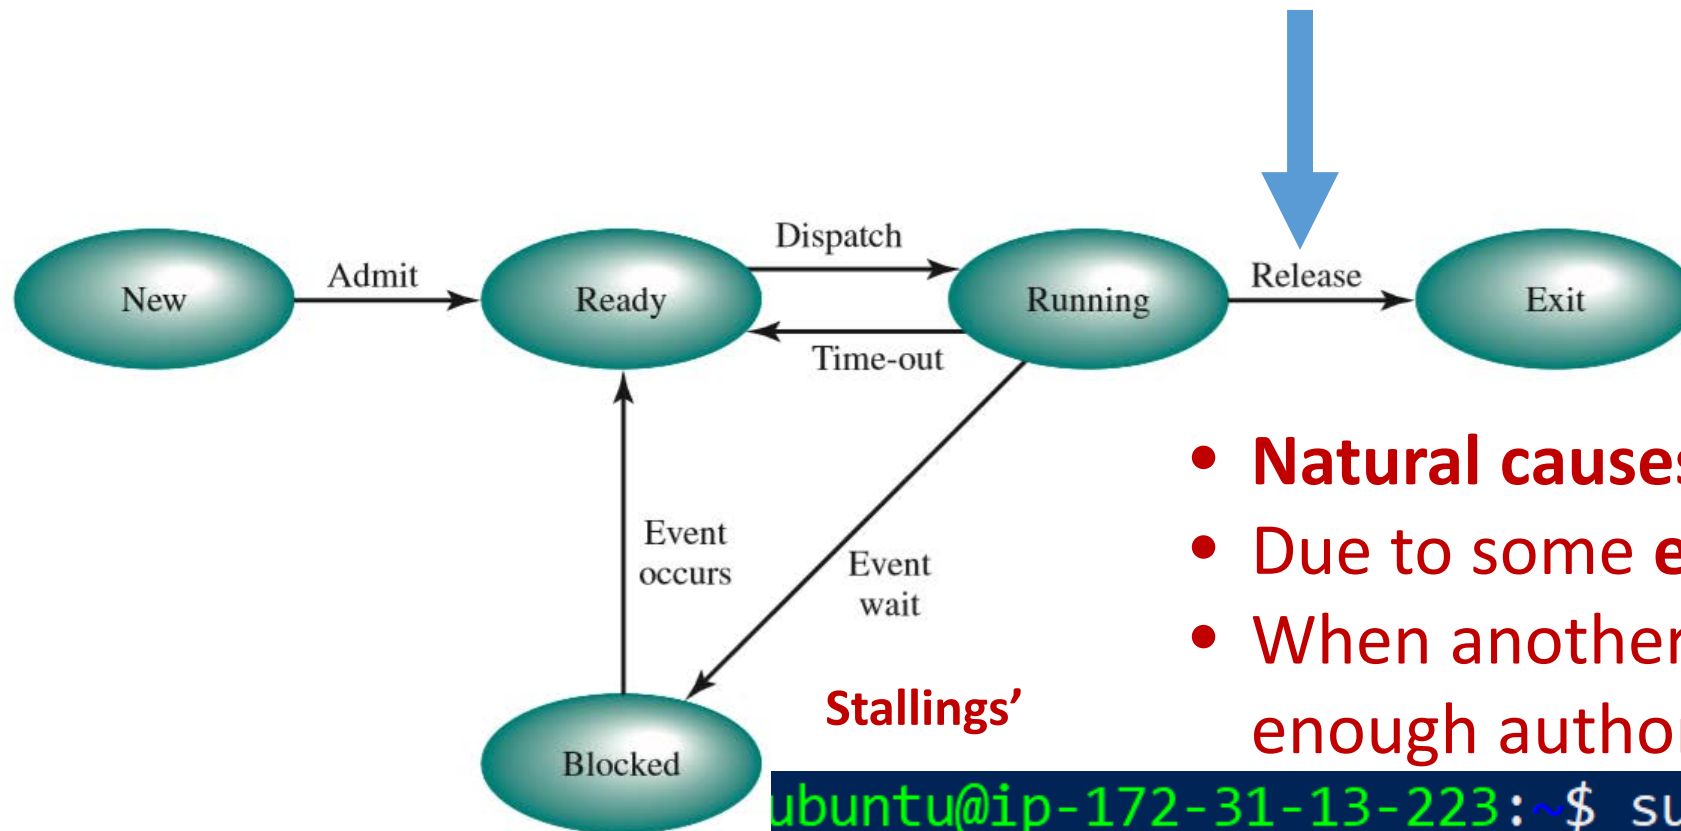
The Life of a process

- A process can get **pre-empted (being robbed of the CPU)** for two reasons:



The Life of a process

- A process can get **terminated (being robbed of its glorious life as process)** for a few reasons:



- **Natural causes**
- **Due to some error/violation**
- **When another process with enough authority kills it.**

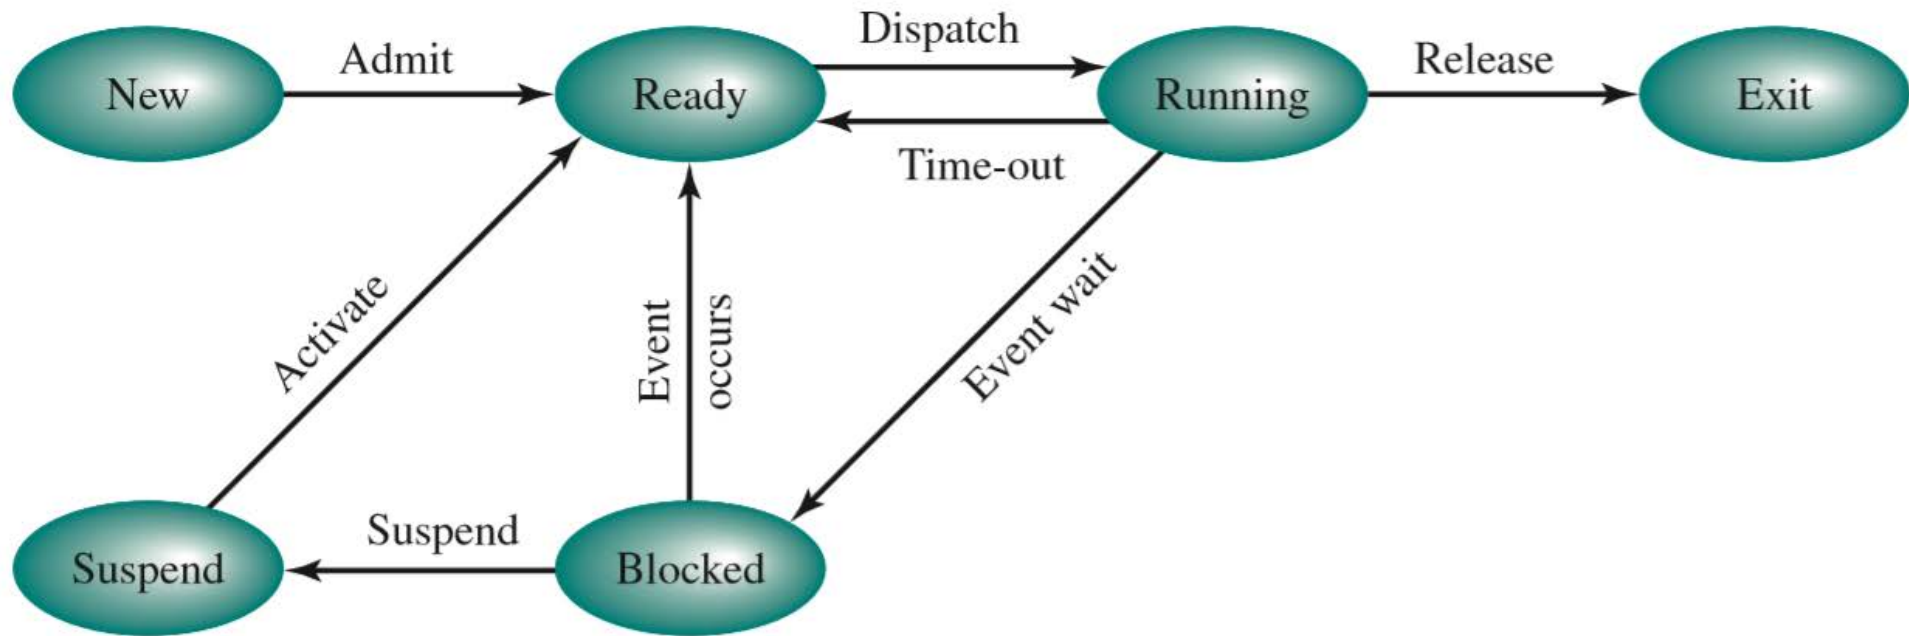
Stallings'


```
ubuntu@ip-172-31-13-223:~$ sudo kill -9 727
```

The Life of a process

Activity (5mins).

Read page 144 from the book and explain this new “suspend” state.



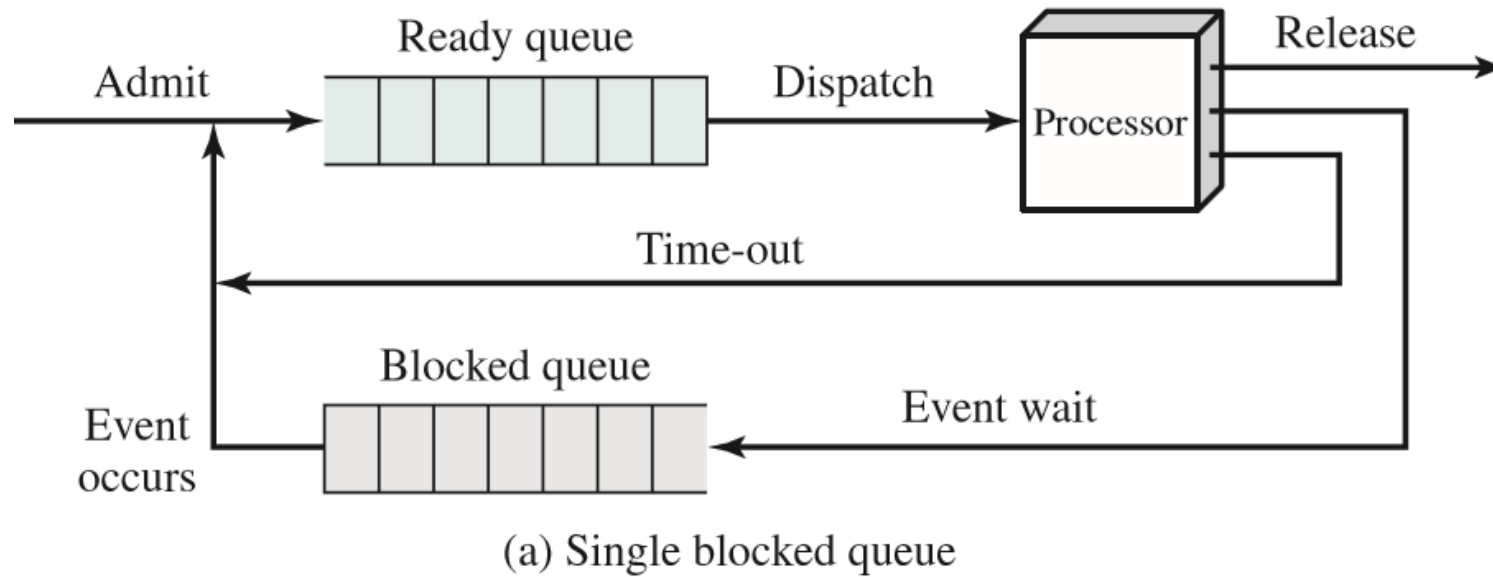
 torvalds Merge git://git.kernel.org/pub/scm/linux/kernel/git/davem

```
70  /* Used in tsk->state: */
71  #define TASK_RUNNING                0x0000
72  #define TASK_INTERRUPTIBLE          0x0001
73  #define TASK_UNINTERRUPTIBLE        0x0002
74  #define __TASK_STOPPED              0x0004
75  #define __TASK_TRACED                0x0008
76  /* Used in tsk->exit_state: */
77  #define EXIT_DEAD                    0x0010
78  #define EXIT_ZOMBIE                  0x0020
79  #define EXIT_TRACE                    (EXIT_ZOMBIE | EXIT_DEAD
```

```
80  /* Used in tsk->state again: */
81  #define TASK_PARKED                  0x0040
82  #define TASK_DEAD                    0x0080
83  #define TASK_WAKEKILL                0x0100
84  #define TASK_WAKING                  0x0200
85  #define TASK_NOLOAD                  0x0400
86  #define TASK_NEW                     0x0800
87  #define TASK_STATE_MAX               0x1000
88  /* Convenience macros for the sake of set_current_state() */
89  #define TASK_KILLABLE                (TASK_WAKEKILL | TASK_DEAD)
90  #define TASK_STOPPED                 (TASK_WAKEKILL | TASK_STOPPED)
91  #define TASK_TRACED                  (TASK_WAKEKILL | TASK_TRACED)
92  #define TASK_IDLE                    (TASK_UNINTERRUPTIBLE | TASK_DEAD)
```


The Life of a process

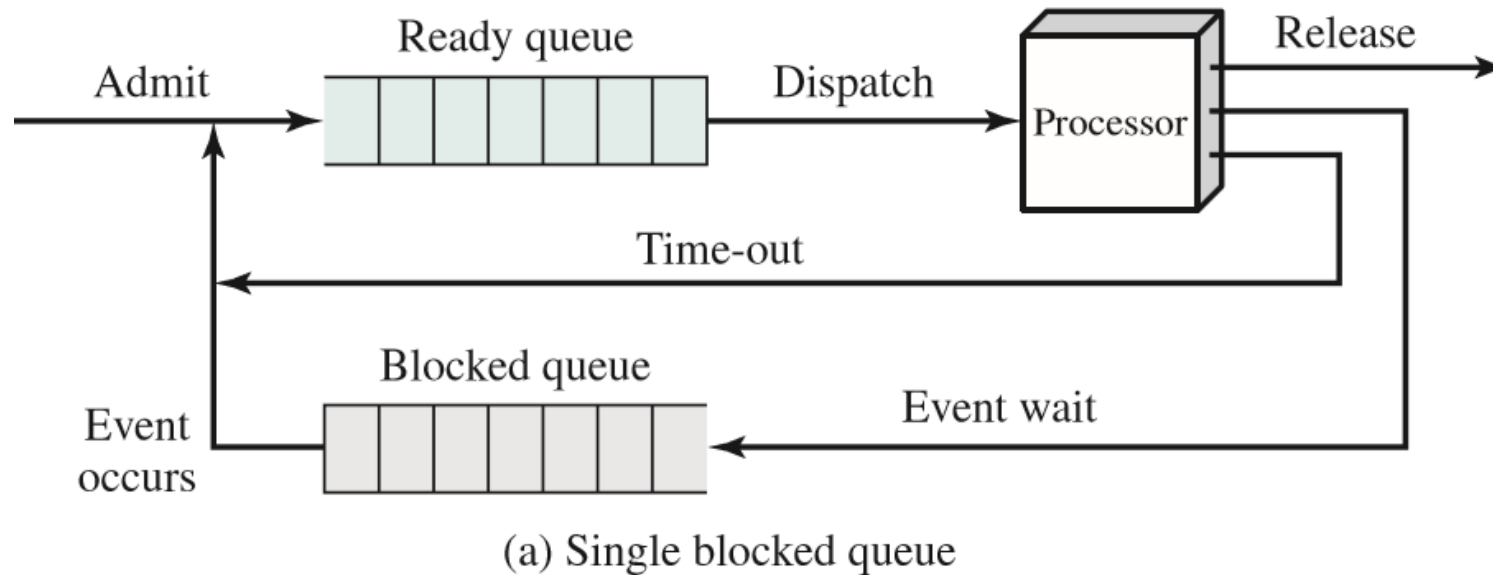
The ready queue and blocked queue



Stallings'

The Life of a process

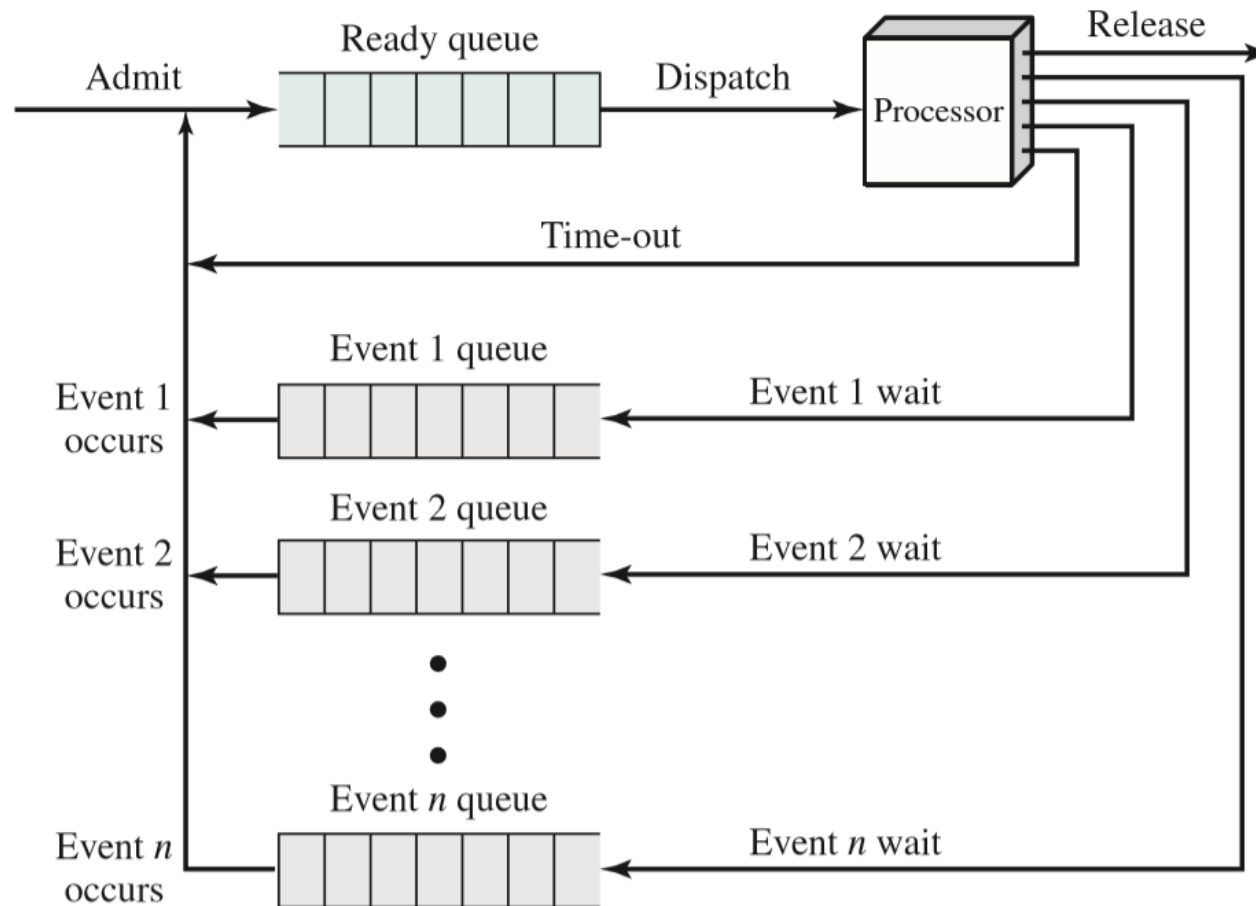
The ready queue and blocked queue



Stallings' **What's the problem with having a single blocked queue?**
Hint: $O(n)$

The Life of a process

The ready queue and blocked queue



Solution:

Multiple Blocked Queues

**When an *event n* occurs
The whole *event n queue*
Is moved to the ready queue**

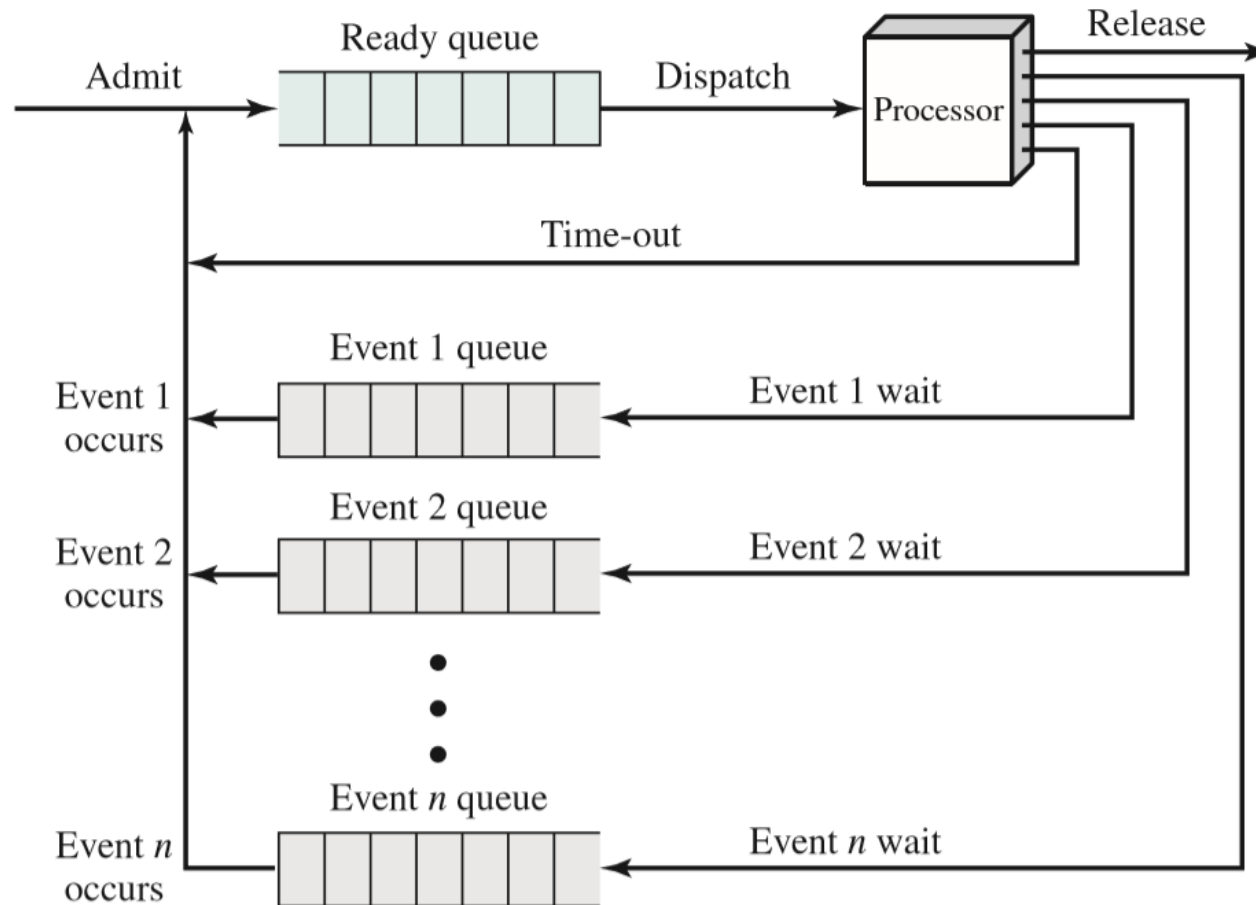
Stallings'

The Life of a process

What exactly can an event be?

That depends on IO drivers
(so many possibilities)

The ready queue and blocked queue



Solution:

Multiple Blocked Queues

When an *event n* occurs

The whole *event n queue*


Is moved to the ready queue

Stallings'

The Life of a process

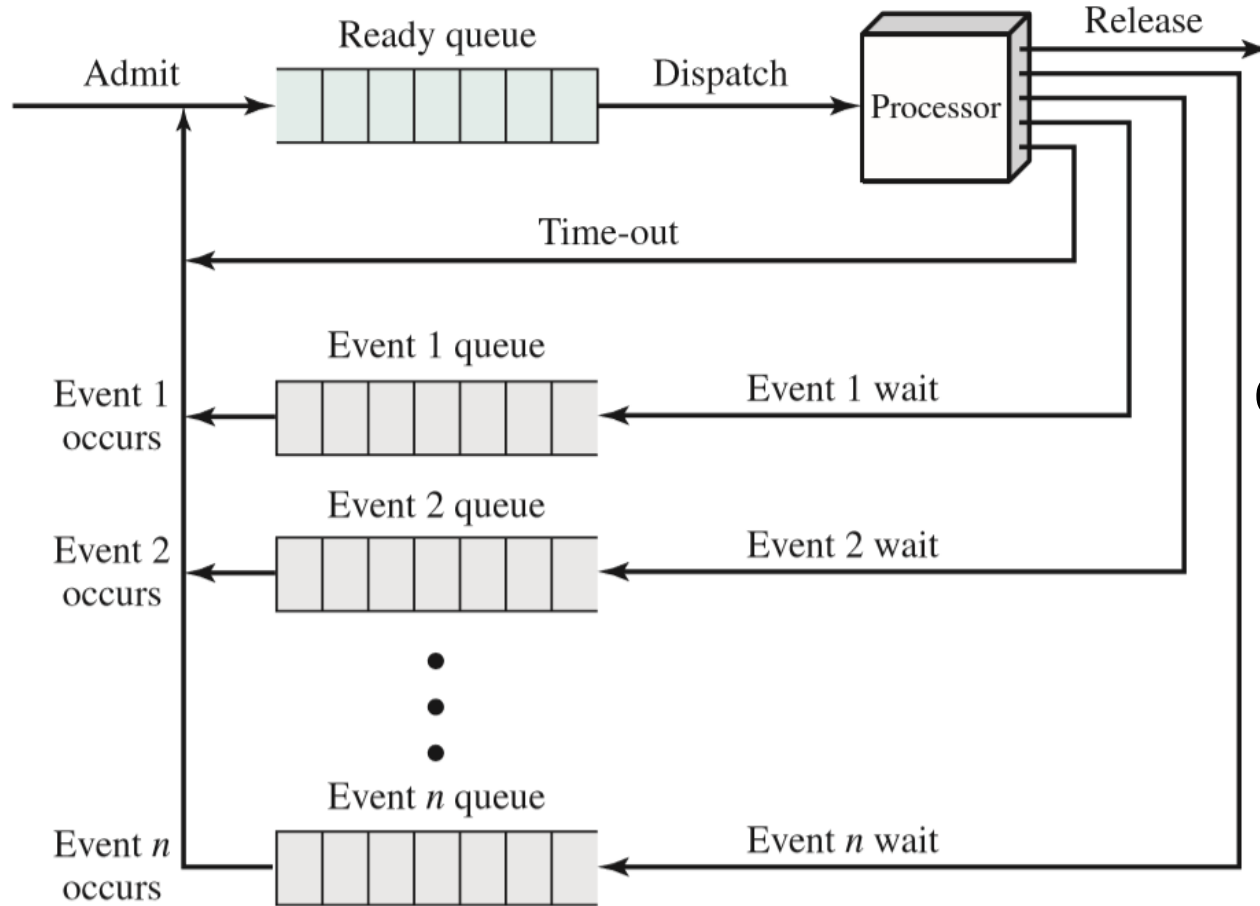
The **ready queue** and **blocked queue**

The Life of a process




Example. Let's say a **user space process** makes a system call to requests a file to the **file system driver**.

The Life of a process




The file system driver will **create a queue** (for that file request) and **put the user process to sleep**

The Life of a process



The file system driver will **then requests the disk driver a sector in disk** where the file contents are stored; and will set up a **second wait queue for that sector in disk.**

The Life of a process




There's some serious magic
Going on here



In this case, the second queue is
**associated with a piece of code (a
callback)**, not really a process.

The Life of a process

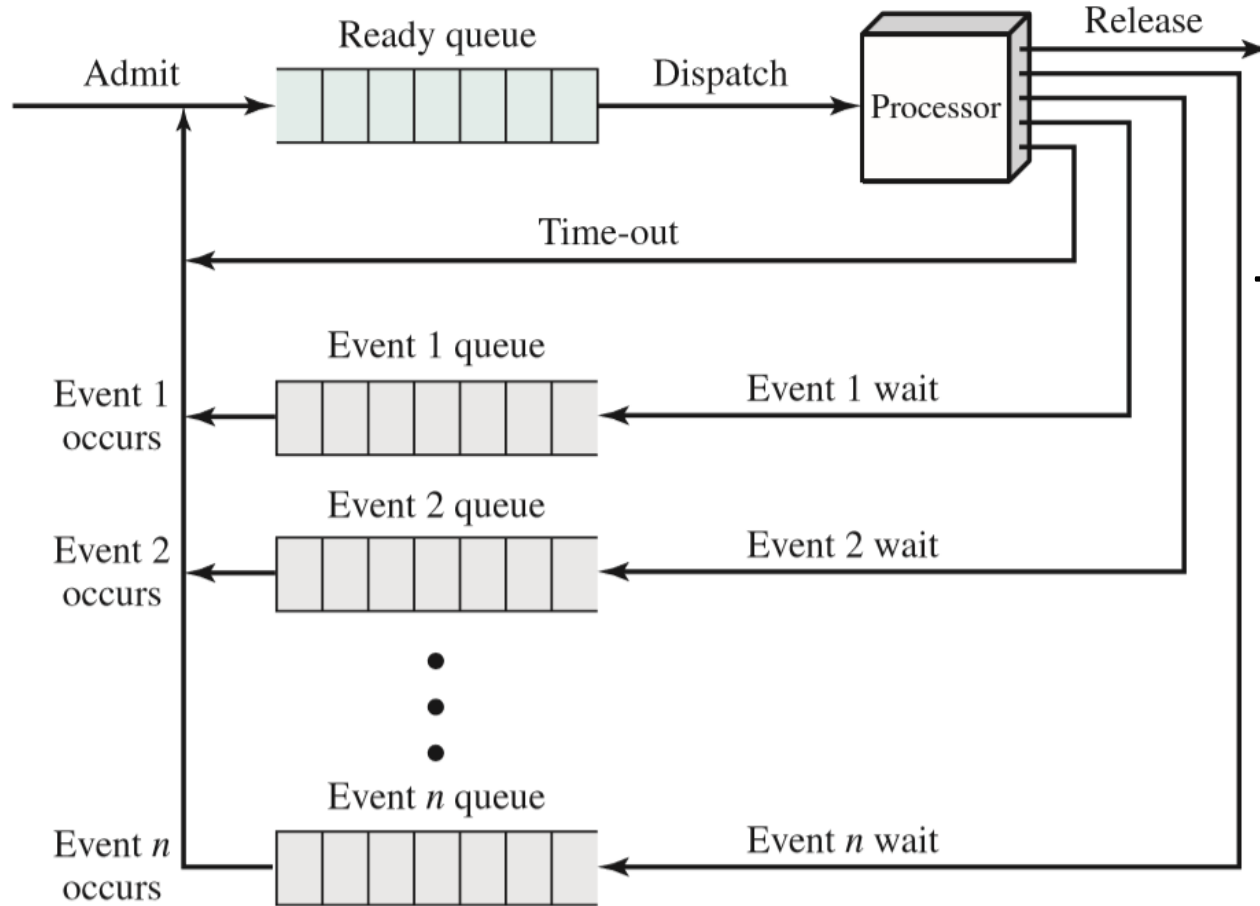


If more requests arrive asking for the **same file**, they can be “placed”/”associated” in the same **wait queue**.



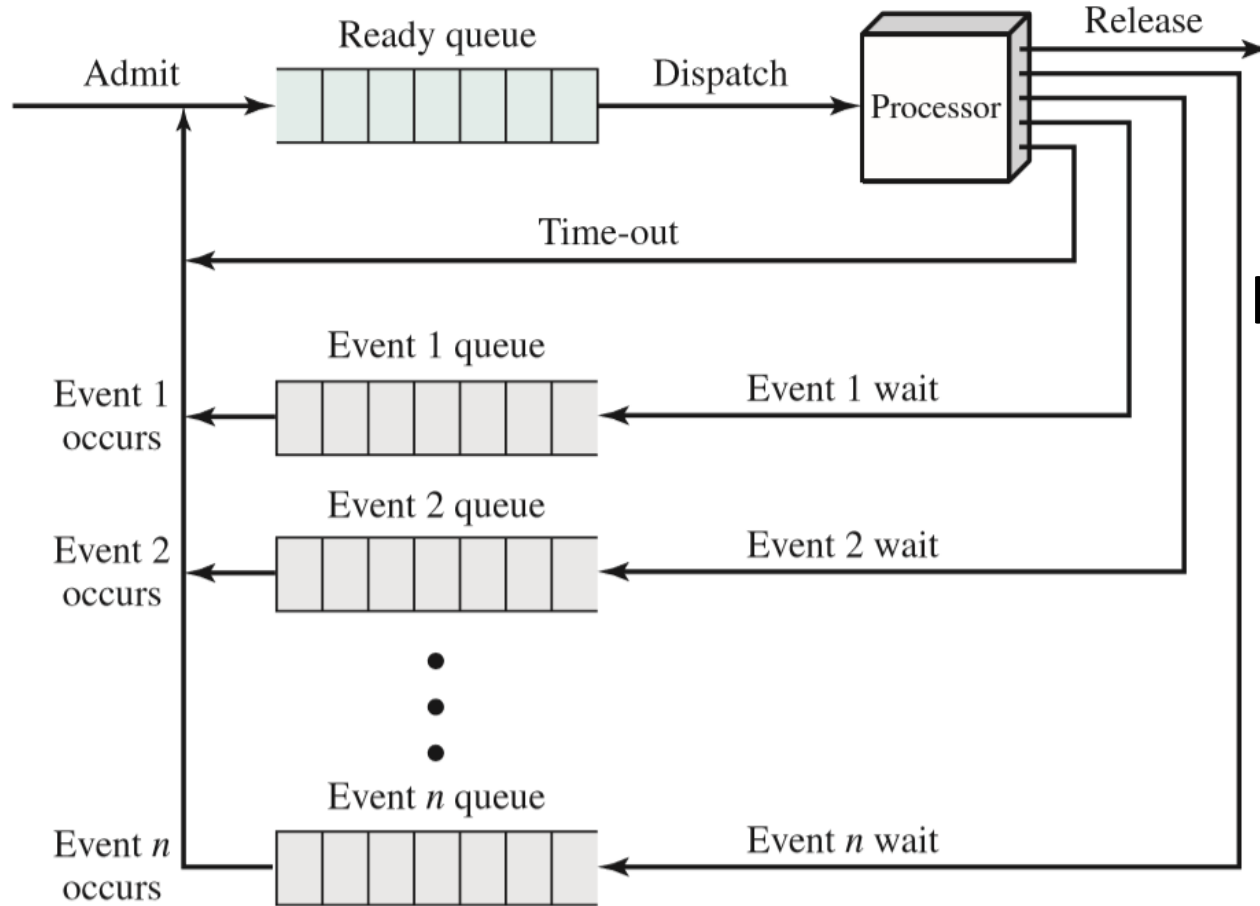
E.g.
Multiple receivers of the same file can
Be handled within the same callback

The Life of a process



The disk driver will then **request the needed sector from disk** and create a **third queue**, and associate a callback with it.

The Life of a process

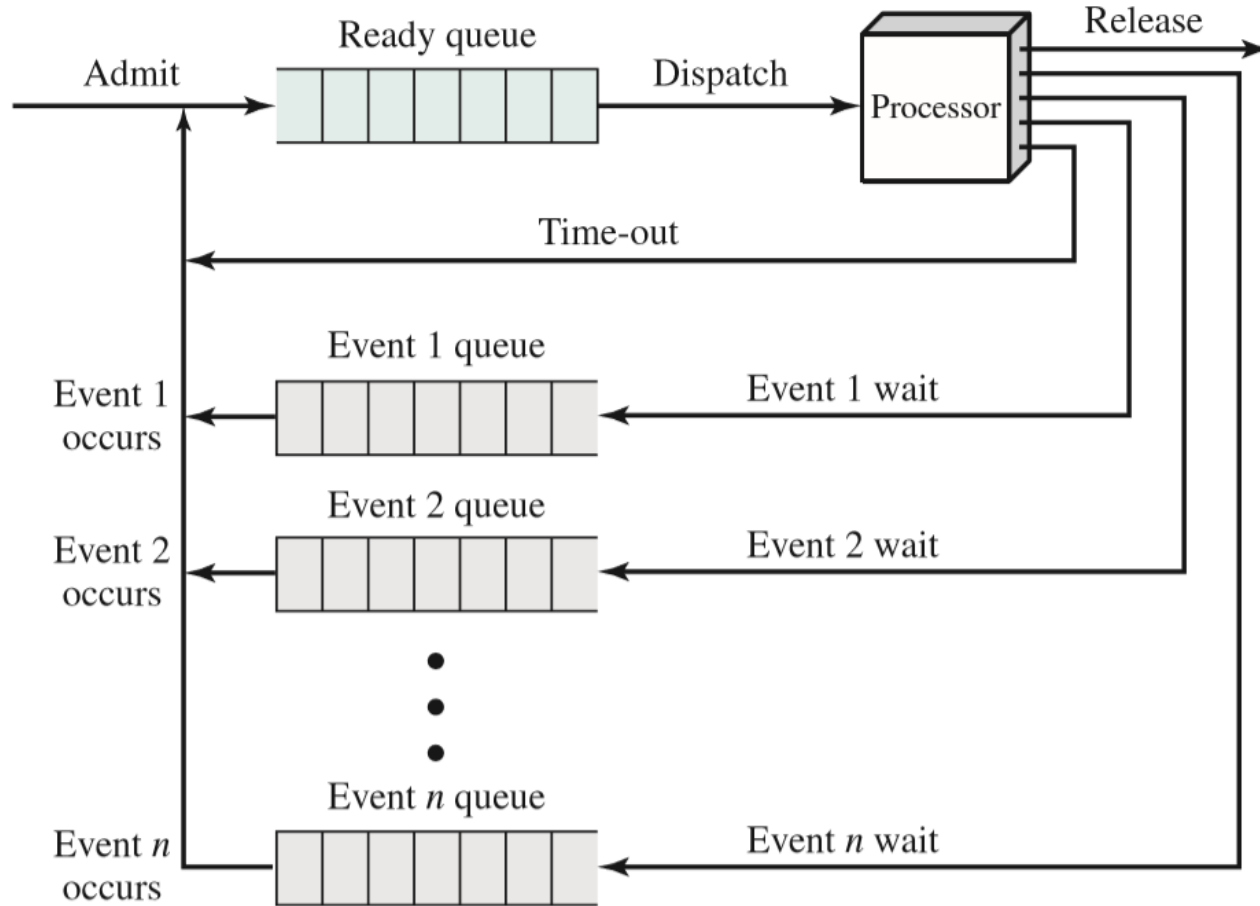


If more requests for the **same sector** arrive. They can be associated with the same queue.



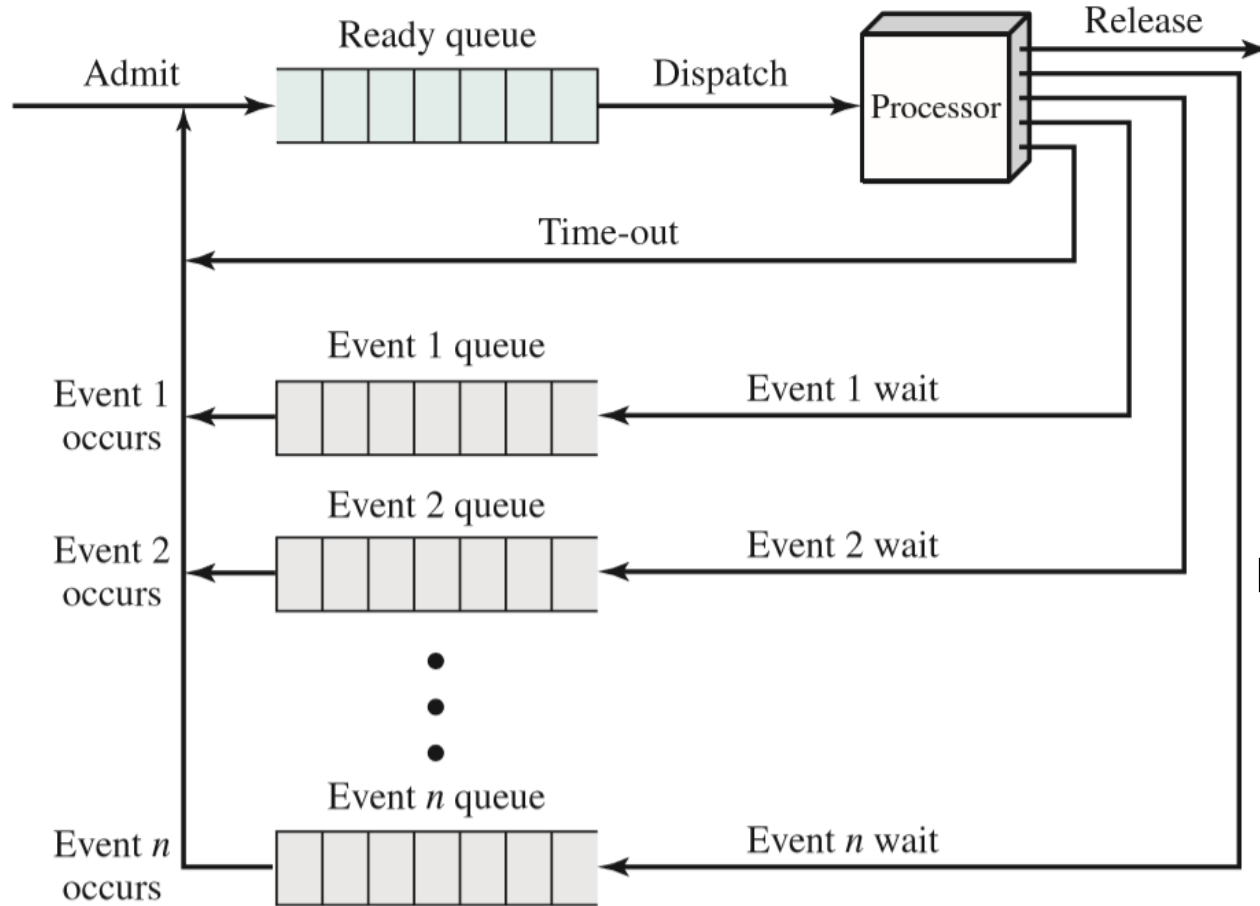
E.g.
**Multiple receivers of the same file can
Be handled within the same callback**

The Life of a process




Eventually when the DMA has finished fetching the sector, it'll trigger a **disk interrupt (part of the disk driver)**.

The Life of a process



In the disk interrupt, the disk driver will **“wake up” the code waiting on the third event queue** (which remember is more disk driver code).

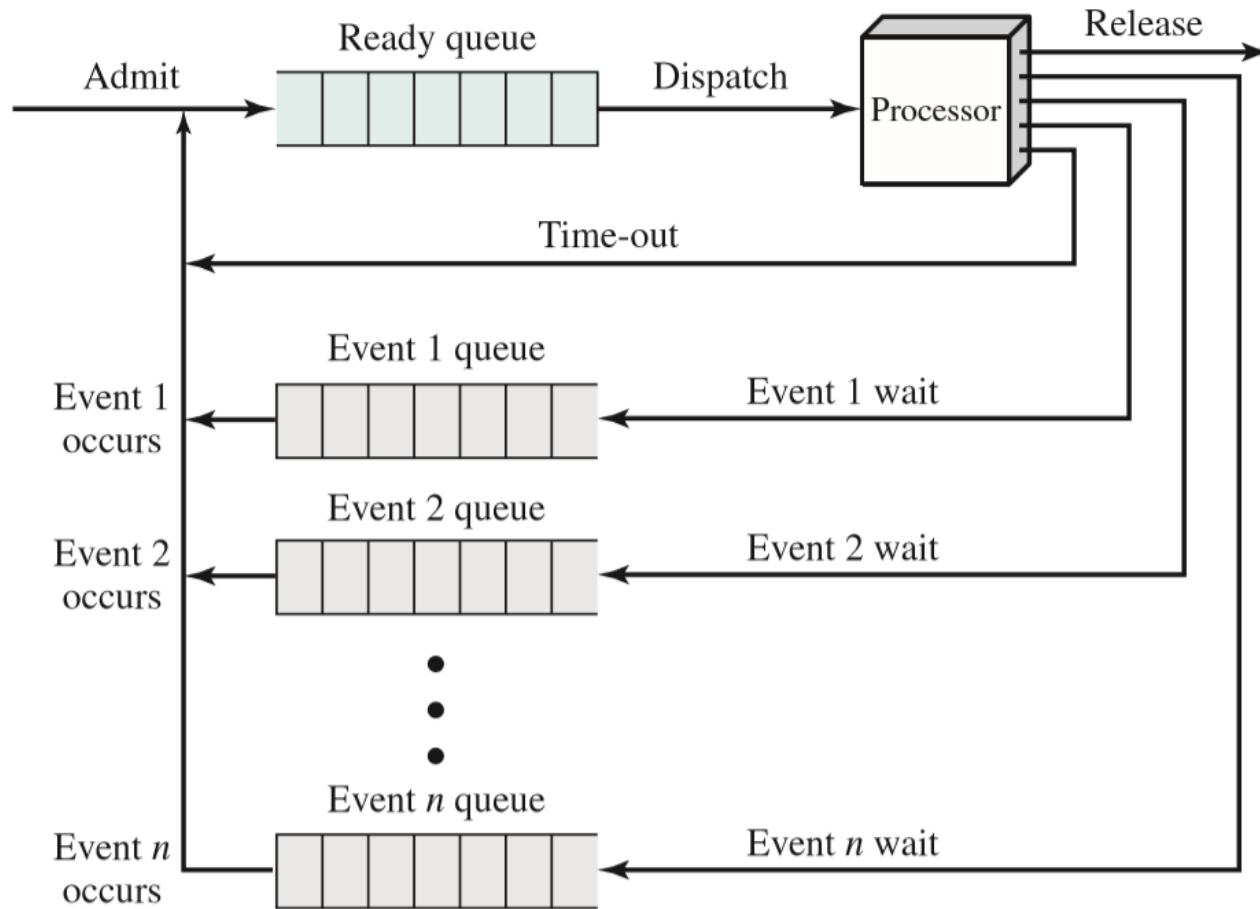
The Life of a process



From the **disk driver** and
Associated with queue 3

This code will **“wake up”** code
waiting on the second queue.


The Life of a process



From the **file driver** and
Associated with queue 2

This code will **“wake up”** the user
process waiting on the first queue.

The Life of a process




Additional notes:

Linux used to have the

*“When an event n occurs
the whole event n queue
is moved to the ready queue”*

semantics, but then in 1999 [a study](#) found that **Linux had some Performance issues.**

The Life of a process



Part of the problem was the *Thundering herd problem*.

Short 5min activity where you find out what it is, and what it has to do with the

“When an event n occurs the whole event n queue is moved to the ready queue”

semantics; and how it was solved.
(hint: exclusive waiting, `prepare_to_wait_exclusive()`)

The Life of a process

Book reference
Section 3.2